

# Cat vs. Dog Image Classification Using Deep Learning in Python

**Carter Susi**

**Dec 02, 2025**

**CAP4613**

**Repo:** <https://github.com/cartersusi/susi-ml-bonus>

```
!pip install kagglehub matplotlib torch torchinfo torchvision tqdm scik
```

```
Requirement already satisfied: kagglehub in /usr/local/lib/python3.12/
Requirement already satisfied: matplotlib in /usr/local/lib/python3.12/
Requirement already satisfied: torch in /usr/local/lib/python3.12/dist-
Collecting torchinfo
    Downloading torchinfo-1.8.0-py3-none-any.whl.metadata (21 kB)
Requirement already satisfied: torchvision in /usr/local/lib/python3.12/
Requirement already satisfied: tqdm in /usr/local/lib/python3.12/dist-
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.
Requirement already satisfied: seaborn in /usr/local/lib/python3.12/di
Requirement already satisfied: packaging in /usr/local/lib/python3.12/
Requirement already satisfied: pyyaml in /usr/local/lib/python3.12/dis
Requirement already satisfied: requests in /usr/local/lib/python3.12/d
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/pyth
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/pyt
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/pyt
Requirement already satisfied: numpy>=1.23 in /usr/local/lib/python3.1
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.12/
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/pyth
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/
Requirement already satisfied: filelock in /usr/local/lib/python3.12/d
Requirement already satisfied: typing-extensions>=4.10.0 in /usr/local/
Requirement already satisfied: setuptools in /usr/local/lib/python3.12
Requirement already satisfied: sympy>=1.13.3 in /usr/local/lib/python3
Requirement already satisfied: networkx>=2.5.1 in /usr/local/lib/pytho
Requirement already satisfied: jinja2 in /usr/local/lib/python3.12/dis
Requirement already satisfied: fsspec>=0.8.5 in /usr/local/lib/python3
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.6.77 in /usr
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.6.77 in /u
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.6.80 in /usr
Requirement already satisfied: nvidia-cudnn-cu12==9.10.2.21 in /usr/lo
Requirement already satisfied: nvidia-cublas-cu12==12.6.4.1 in /usr/lo
Requirement already satisfied: nvidia-cufft-cu12==11.3.0.4 in /usr/loc
Requirement already satisfied: nvidia-curand-cu12==10.3.7.77 in /usr/l
Requirement already satisfied: nvidia-cusolver-cu12==11.7.1.2 in /usr/
Requirement already satisfied: nvidia-cusparse-cu12==12.5.4.2 in /usr/
```

```
Requirement already satisfied: nvidia-cusparselt-cu12==0.7.1 in /usr/local/lib/python3.11/site-packages/nvidia/cusparselt/_cusparselt.cpython-311.pd
Requirement already satisfied: nvidia-nccl-cu12==2.27.5 in /usr/local/lib/python3.11/site-packages/nvidia/nccl/_nccl.cpython-311.pd
Requirement already satisfied: nvidia-nvshmem-cu12==3.3.20 in /usr/local/lib/python3.11/site-packages/nvidia/nvshmem/_nvshmem.cpython-311.pd
Requirement already satisfied: nvidia-nvtx-cu12==12.6.77 in /usr/local/lib/python3.11/site-packages/nvidia/nvtx/_nvtx.cpython-311.pd
Requirement already satisfied: nvidia-nvjitlink-cu12==12.6.85 in /usr/local/lib/python3.11/site-packages/nvidia/nvjitlink/_nvjitlink.cpython-311.pd
Requirement already satisfied: nvidia-cufile-cu12==1.11.1.6 in /usr/local/lib/python3.11/site-packages/nvidia/cufile/_cufile.cpython-311.pd
Requirement already satisfied: triton==3.5.0 in /usr/local/lib/python3.11/site-packages/triton/_triton.cpython-311.pd
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.11/site-packages/scipy/_scipy.cpython-311.pd
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.11/site-packages/joblib/_joblib.cpython-311.pd
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.11/site-packages/threadpoolctl/_threadpoolctl.cpython-311.pd
Requirement already satisfied: pandas>=1.2 in /usr/local/lib/python3.11/site-packages/pandas/_pandas.cpython-311.pd
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/site-packages/pytz/_pytz.cpython-311.pd
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/site-packages/tzdata/_tzdata.cpython-311.pd
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/site-packages/six/_six.cpython-312.pd
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.11/site-packages/mpmath/_mpmath.cpython-311.pd
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.11/site-packages/markupsafe/_markupsafe.cpython-311.pd
Requirement already satisfied: charset_normalizer<4,>=2 in /usr/local/lib/python3.11/site-packages/charset_normalizer/_charset_normalizer.cpython-311.pd
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/site-packages/idna/_idna.cpython-311.pd
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/site-packages/urllib3/_urllib3.cpython-311.pd
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/site-packages/certifi/_certifi.cpython-311.pd
  Downloading torchinfo-1.8.0-py3-none-any.whl (23 kB)
  Installing collected packages: torchinfo
```

## ▼ Set Config

```
import os
import shutil

import torch
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

import warnings
from PIL import Image, ImageFile

warnings.filterwarnings("ignore", category=UserWarning, module="PIL")
warnings.filterwarnings("ignore", ".*Truncated File Read.*")
ImageFile.LOAD_TRUNCATED_IMAGES = True

cwd = os.getcwd()
home = os.path.expanduser("~/")

conf = {
    "models_dir": f"{ cwd }/models",
    "kaggle_creds": f"{ home }/.kaggle/kaggle.json",
    "model_path": f"/{ cwd }/models/cat_dog_cnn_25.pth"
}

def copy_creds(from_path: str, to_path: str) -> bool:
    try:
        os.makedirs(os.path.dirname(conf["kaggle_creds"])), exist_ok=True
        . . . . .
```

```
shutil.copy(from_path, to_path)
os.chmod(to_path, 0o600)
except Exception as e:
    print("error copying kaggle credential file", e)

if not os.path.exists(conf["kaggle_creds"]):
    try:
        from google.colab import files
        print("Please upload your 'kaggle.json' credential file\n")
        files.upload()
        copy_creds("kaggle.json", conf["kaggle_creds"])
    except ImportError:
        if not os.path.exists("kaggle.json"):
            raise Exception(f"kaggle.json is not in current directory or {ci}
copy_creds("kaggle.json", conf["kaggle_creds"])

try:
    import google.colab

    image_size = (64, 64)
    print(f"In Colab, using image size: '{image_size[0]}x{image_size[1]}')
except ImportError:
    vram = torch.cuda.get_device_properties(0).total_memory / (1024**3)
    if vram < 10.0: # Needs ~8.5 Gb for 256x256
        image_size = (128, 128)
    else:
        image_size = (256, 256)

    print(
        f"Not in Colab with {vram} GB VRAM, default image size: '{image_s:
    )
```

```
Please upload your 'kaggle.json' credential file
 kaggle.json
kaggle.json(application/json) - 66 bytes, last modified: n/a - 100% done
Saving kaggle.json to kaggle.json
In Colab, using image size: '64x64px'
```

## ▼ Download Cat Vs Dog Dataset

```
import kagglehub

def _find_image_folder(root_path: str) -> str:
    for dirpath, dirnames, _ in os.walk(root_path):
        if "Cat" in dirnames and "Dog" in dirnames:
            return dirpath
    return ""
```

```
dataset_path = _find_image_folder()
kagglehub.dataset_download("karakaggle/kaggle-cat-vs-dog-dataset")
)

print(f"Kaggle Dataset at: {dataset_path}")
```

```
Using Colab cache for faster access to the 'kaggle-cat-vs-dog-dataset'
Kaggle Dataset at: /kaggle/input/kaggle-cat-vs-dog-dataset/kagglecatsanddogs_3329image
```

## Load Dataset

```
import random

import matplotlib.pyplot as plt
import torch
from torch.utils.data import DataLoader, Subset, random_split
from torchvision import datasets, transforms
from torchvision.transforms import v2 as T

class Dataset:
    def __init__(self, fpath):
        self.data_path = fpath
        self.dataset = datasets.ImageFolder(root=fpath)

    # Training image transforms (70% training)
    self.train_transform = transforms.Compose(
        [
            transforms.Resize(image_size), # Resize all images to 150x150
            transforms.RandomHorizontalFlip(), # Apply basic data augmentation
            transforms.RandomRotation(
                15
            ), # Apply basic data augmentation (rotation)
            transforms.ToTensor(), # Normalize pixel values to [0, 1] range
            transforms.RandomApply(
                [T.GaussianNoise(mean=0.0, sigma=0.20, clip=True)],
                p=0.5, # Apply max-noise of 0.20 with 50% probability - randomly
            ),
        ]
    )

    # Validation image transforms (15% validation)
    self.validation_transform = transforms.Compose(
        [
            transforms.Resize(image_size),
            transforms.ToTensor(),
        ]
    )


```

```
# Test image transforms (15% val)
self.test_transform = self.validation_transform

@property
def data_size(self) -> int:
    return len(self.dataset)

@property
def train_size(self) -> int:
    return int(0.7 * self.data_size)

@property
def validation_size(self) -> int:
    return int(0.15 * self.data_size)

@property
def test_size(self) -> int:
    return int(self.data_size - self.train_size - self.validation_size)

def preprocess(self):
    if not self.dataset:
        raise Exception("Dataset must be loaded before preprocessing.")

    train_start, validation_start, test_start = random_split(
        self.dataset, [self.train_size, self.validation_size, self.test_size]
    )

    self.train_data = Subset(
        datasets.ImageFolder(self.data_path, transform=self.train_transform),
        train_start.indices,
    )
    self.val_data = Subset(
        datasets.ImageFolder(self.data_path, transform=self.validation_transform),
        validation_start.indices,
    )
    self.test_data = Subset(
        datasets.ImageFolder(self.data_path, transform=self.test_transform),
        test_start.indices,
    )

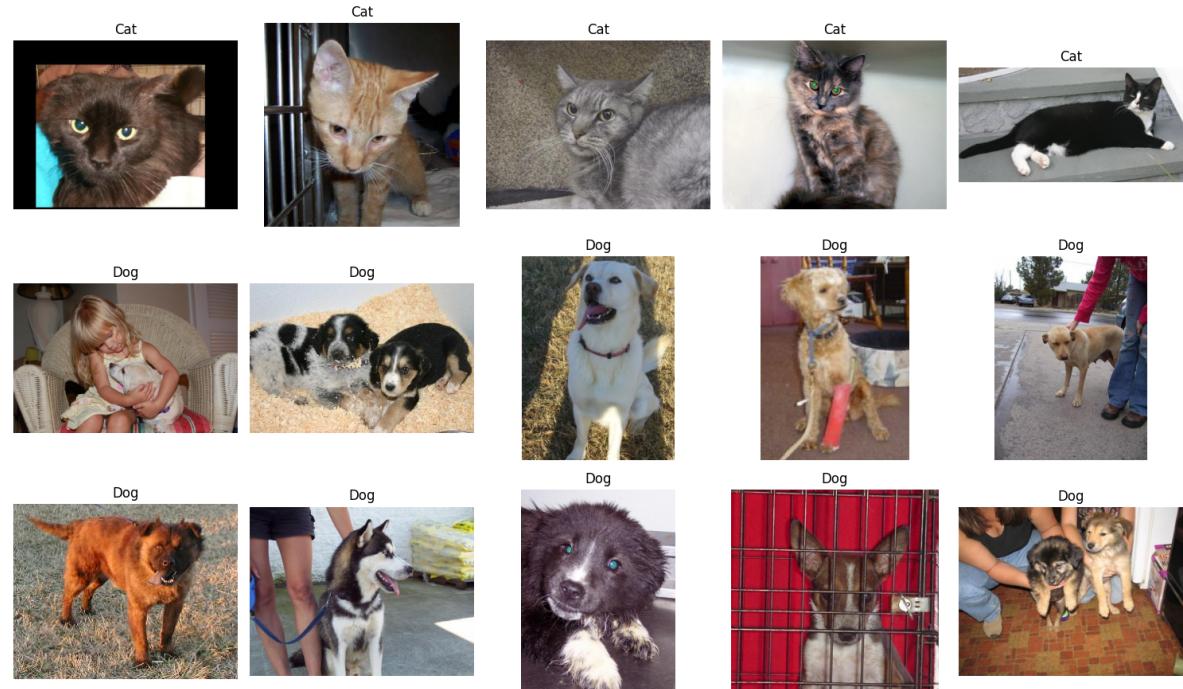
def show_samples(self, train_data=False):
    cats = [
        i for i, (path, _) in enumerate(self.dataset.samples) if "Cat" in path
    ]
    dogs = [
        i for i, (path, _) in enumerate(self.dataset.samples) if "Dog" in path
    ]

    sample_cats = random.sample(cats, min(10, len(cats)))
    sample_dogs = random.sample(dogs, min(10, len(dogs)))
```

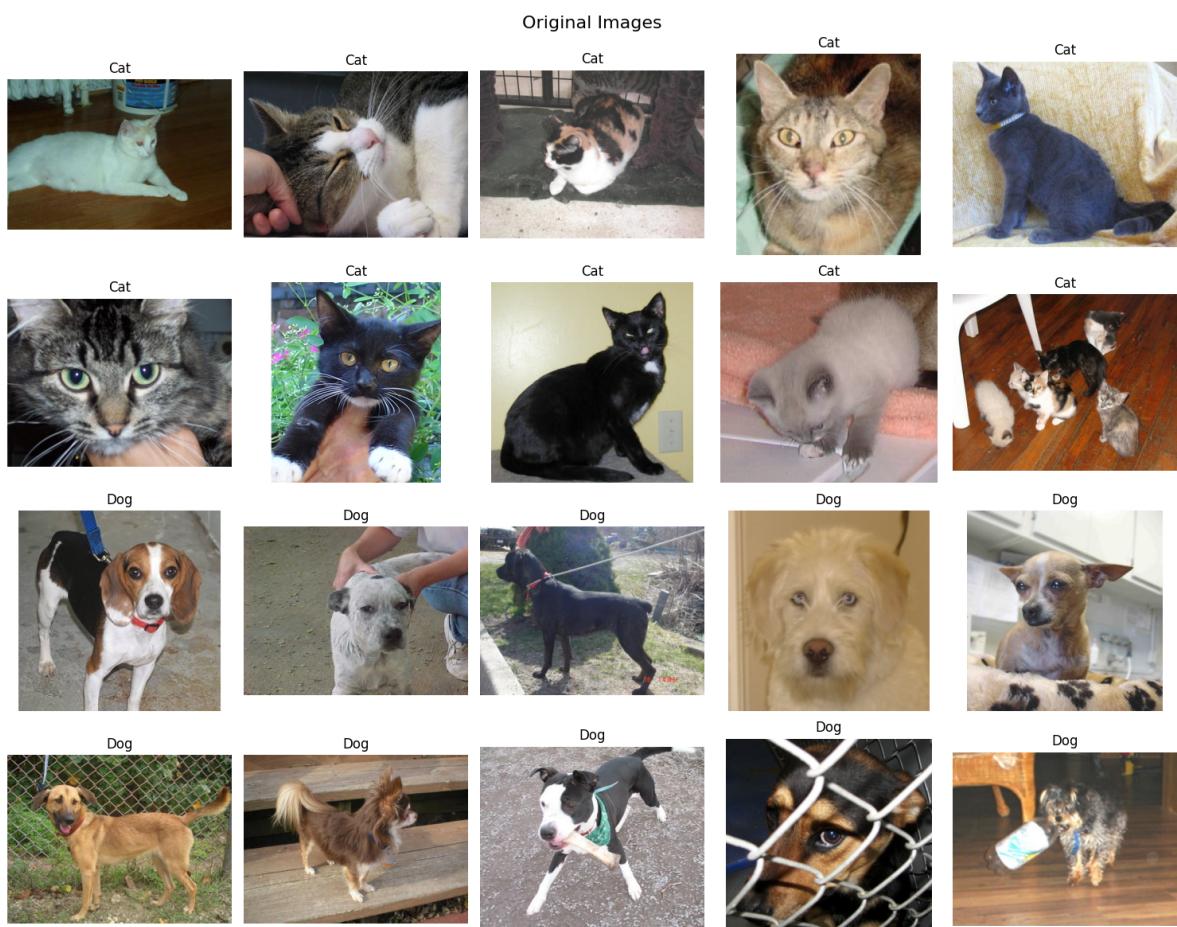
```
sample_cats = random.sample(cats, min(10, len(cats)),  
title = "Augmented Images" if train_data else "Original Images"  
fig, axes = plt.subplots(4, 5, figsize=(15, 12))  
fig.suptitle(title, fontsize=16)  
  
for idx, ax in enumerate(axes.flat):  
    img_idx = sample_cats[idx % 10] if idx < 10 else sample_dogs[idx % 10]  
    img, _ = self.dataset[img_idx]  
  
    if train_data:  
        img = self.train_transform(img).permute(1, 2, 0).numpy()  
  
    ax.imshow(img)  
    ax.set_title("Cat" if idx < 10 else "Dog")  
    ax.axis("off")  
  
plt.tight_layout()  
plt.show()  
  
def augment(self, datatype: str):  
    torch.manual_seed(42)  
  
    if datatype == "train":  
        return DataLoader(  
            self.train_data, batch_size=32, shuffle=True, num_workers=os.cpu_count()  
        )  
    else:  
        return DataLoader(  
            self.val_data, batch_size=32, shuffle=True, num_workers=os.cpu_count()  
        )  
  
dataset = Dataset(dataset_path)  
  
print("Dataset samples before processing.\n")  
dataset.show_samples()  
  
dataset.preprocess()  
  
print("\nDataset samples after processing.\n")  
dataset.show_samples()
```

Dataset samples before processing.





Dataset samples after processing.





## ▼ Create Model

```
import torch.nn as nn
import torchinfo

class CatDogCNN(nn.Module):
    def __init__(self):
        super().__init__()

        self.conv_layer_1 = nn.Sequential(
            nn.Conv2d(3, 64, 3, padding=1),
            nn.ReLU(),
            nn.BatchNorm2d(64),
            nn.MaxPool2d(2),
        )

        self.conv_layer_2 = nn.Sequential(
            nn.Conv2d(64, 512, 3, padding=1),
            nn.ReLU(),
            nn.BatchNorm2d(512),
            nn.MaxPool2d(2),
        )

        self.conv_layer_3 = nn.Sequential(
            nn.Conv2d(512, 512, kernel_size=3, padding=1),
            nn.ReLU(),
            nn.BatchNorm2d(512),
            nn.MaxPool2d(2)
        )
```

```
        nn.MaxPool2d(2),
    )

    self.adaptive_pool = nn.AdaptiveAvgPool2d((2, 2))

    self.classifier = nn.Sequential(
        nn.Flatten(), nn.Linear(in_features=512 * 2 * 2, out_features=2
    )

def forward(self, x: torch.Tensor) -> torch.Tensor:
    x = self.conv_layer_1(x)
    x = self.conv_layer_2(x)
    x = self.conv_layer_3(x)
    x = self.conv_layer_3(x)
    x = self.conv_layer_3(x)
    x = self.conv_layer_3(x)
    x = self.adaptive_pool(x)
    x = self.classifier(x)
    return x

def summary(self) -> None:
    torchinfo.summary(self, input_size=[1, 3, image_size[0], image_si:

model = CatDogCNN()
model.summary()
model.to(device)
```

```
CatDogCNN(
    (conv_layer_1): Sequential(
        (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
        (1): ReLU()
        (2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    )
    (conv_layer_2): Sequential(
        (0): Conv2d(64, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
        (1): ReLU()
        (2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    )
    (conv_layer_3): Sequential(
        (0): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
        (1): ReLU()
        (2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
```

```
(3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
)
(adaptive_pool): AdaptiveAvgPool2d(output_size=(2, 2))
(classifier): Sequential(
    (0): Flatten(start_dim=1, end_dim=-1)
    (1): Linear(in_features=2048, out_features=2, bias=True)
)
)
```

## ▼ Train Model

```
from timeit import default_timer as timer
from typing import Any, Literal, overload

from tqdm import tqdm

train_dataloader_augmented = dataset.augment("train")
val_dataloader_augmented = dataset.augment("val")

loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(params=model.parameters(), lr=1e-3)

def train() -> dict[str, list[Any]]:
    torch.manual_seed(42)
    torch.cuda.manual_seed(42)
    os.makedirs(conf["models_dir"], exist_ok=True)

    start_time = timer()

    model_results = _train(epochs=25)

    end_time = timer()
    print(f"Total training time: {end_time - start_time:.3f} seconds")

    return model_results

def _train(epochs: int = 25) -> dict[str, list[Any]]:
    results = {"train_loss": [], "train_acc": [], "val_loss": [], "val_
```

```
for epoch in tqdm(range(epochs)):
    train_loss, train_acc = _train_step()
```

```
    val_loss, val_acc = _val_step(return_preds=False)
```

```
    tqdm.write(
        f"Epoch: {epoch + 1} | "
        f"train_loss: {train_loss:.4f} | "
```

```
f"train_acc: {train_acc:.4f} | "
f"val_loss: {val_loss:.4f} | "
f"val_acc: {val_acc:.4f}"
)

results["train_loss"].append(train_loss)
results["train_acc"].append(train_acc)
results["val_loss"].append(val_loss)
results["val_acc"].append(val_acc)

torch.save(
    model.state_dict(), os.path.join(conf["models_dir"], f"cat_dog_"
)

return results

def _train_step() -> tuple[float, float]:
    model.train()

    train_loss, train_acc = 0, 0
    for _, (X, y) in enumerate(train_dataloader_augmented):
        X, y = X.to(device), y.to(device)

        # 1. Forward pass
        y_pred = model(X)

        # 2. Calculate and accumulate loss
        loss = loss_fn(y_pred, y)
        train_loss += loss.item()

        # 3. Optimizer zero grad
        optimizer.zero_grad()

        # 4. Loss backward
        loss.backward()

        # 5. Optimizer step
        optimizer.step()

        y_pred_class = torch.argmax(torch.softmax(y_pred, dim=1), dim=1)
        train_acc += (y_pred_class == y).sum().item() / len(y_pred)

    train_loss = train_loss / len(train_dataloader_augmented)
    train_acc = train_acc / len(train_dataloader_augmented)
    return train_loss, train_acc

@overload
def _val_step(return_preds: Literal[False] = False) -> tuple[float, f
@overload
```

```
def _val_step(return_preds: Literal[True]) -> tuple[float, float, list]:
    def _val_step(return_preds: bool = False) -> tuple[float, float] | tuple[None, None]:
        model.eval()
        val_loss, val_acc = 0, 0
        all_preds, all_labels = [], []

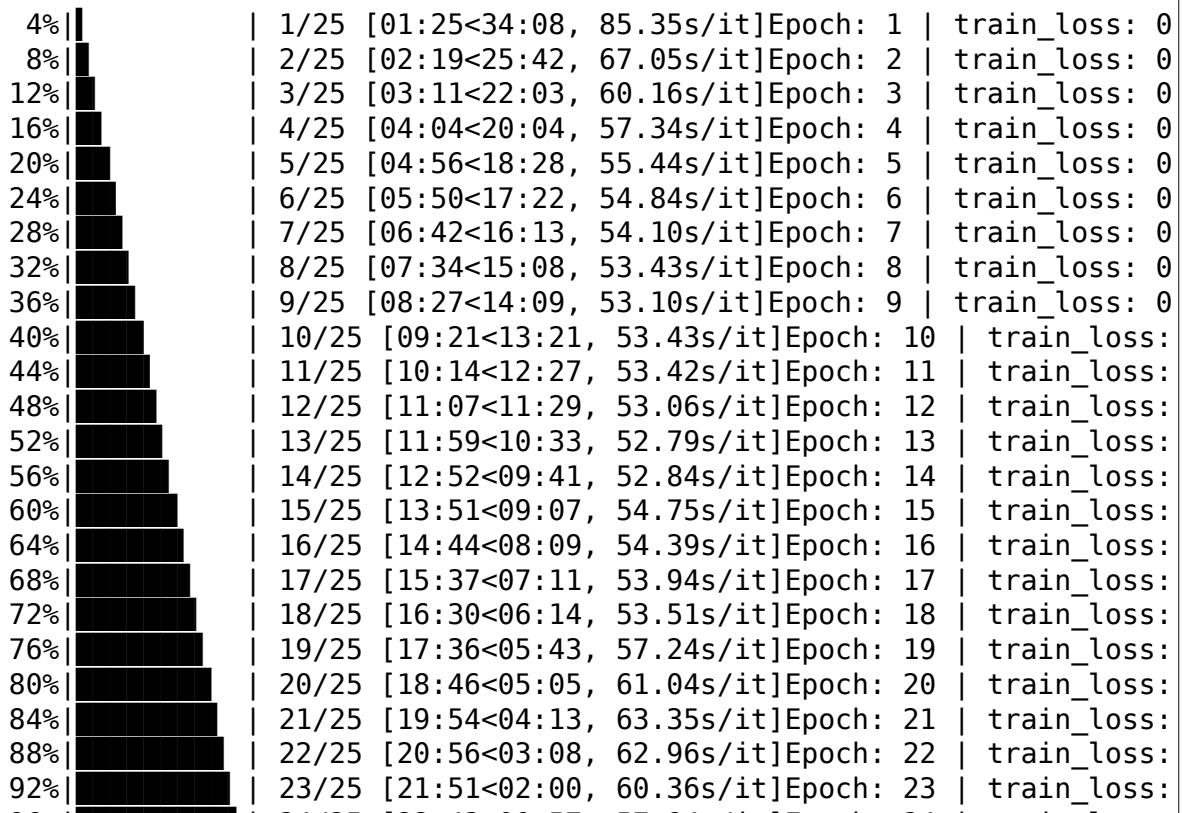
        with torch.inference_mode():
            for _, (X, y) in enumerate(val_dataloader_augmented):
                X, y = X.to(device), y.to(device)
                val_pred_logits = model(X)
                loss = loss_fn(val_pred_logits, y)
                val_loss += loss.item()
                val_pred_labels = val_pred_logits.argmax(dim=1)
                val_acc += (val_pred_labels == y).sum().item() / len(val_pred_labels)

        if return_preds:
            all_preds.extend(val_pred_labels.cpu().numpy())
            all_labels.extend(y.cpu().numpy())

        val_loss = val_loss / len(val_dataloader_augmented)
        val_acc = val_acc / len(val_dataloader_augmented)

    if return_preds:
        return val_loss, val_acc, all_preds, all_labels
    return val_loss, val_acc
```

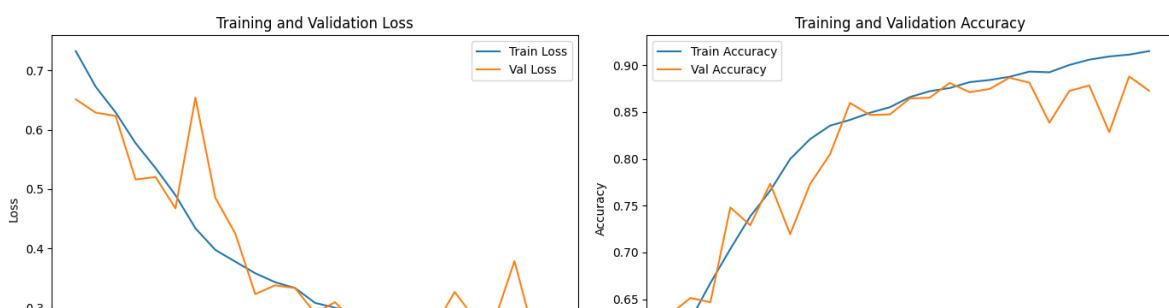
```
res = train()
```

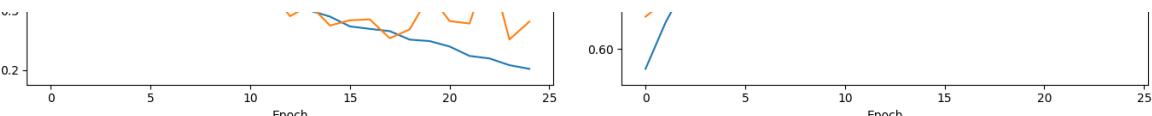


```
96%|██████████| 24/25 [22:43<00:5/, 5/.84s/it]Epoch: 24 | train_loss:  
100%|██████████| 25/25 [23:35<00:00, 56.61s/it]Epoch: 25 | train_loss:  
Total training time: 1415.255 seconds
```

## ▼ Visualizations

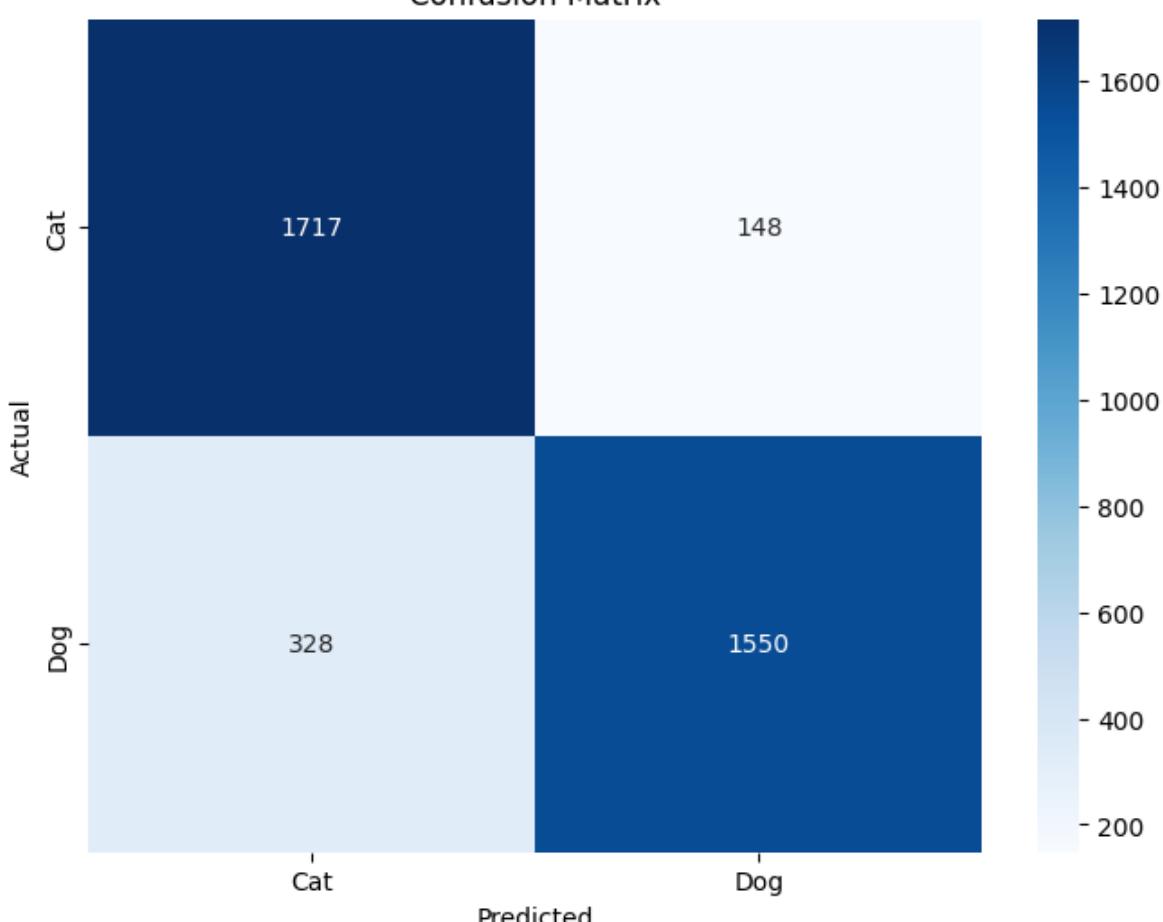
```
from sklearn.metrics import confusion_matrix  
import seaborn as sns  
  
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(14, 5))  
ax1.plot(res['train_loss'], label='Train Loss')  
ax1.plot(res['val_loss'], label='Val Loss')  
ax1.set_xlabel('Epoch')  
ax1.set_ylabel('Loss')  
ax1.set_title('Training and Validation Loss')  
ax1.legend()  
  
ax2.plot(res['train_acc'], label='Train Accuracy')  
ax2.plot(res['val_acc'], label='Val Accuracy')  
ax2.set_xlabel('Epoch')  
ax2.set_ylabel('Accuracy')  
ax2.set_title('Training and Validation Accuracy')  
ax2.legend()  
plt.tight_layout()  
plt.show()  
  
model.eval()  
_, val_acc, val_preds, val_labels = _val_step(return_preds=True)  
  
print(f'\nFinal Validation Accuracy: {val_acc:.4f}')  
  
cm = confusion_matrix(val_labels, val_preds)  
plt.figure(figsize=(8, 6))  
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['Cat'  
plt.xlabel('Predicted')  
plt.ylabel('Actual')  
plt.title('Confusion Matrix')  
plt.show()
```





Final Validation Accuracy: 0.8728

Confusion Matrix



## Model Inference - Upload Image

```
#@title Model Inference - Upload Image

model.eval()

transform = transforms.Compose(
    [
        transforms.Resize(image_size),
        transforms.ToTensor(),
    ]
)

uploaded = files.upload()

image_path = list(uploaded.keys())[0]
image = Image.open(image_path).convert("RGB")
image_tensor = transform(image)
image_tensor = image_tensor.unsqueeze(0).to(device)

with torch.inference_mode():
    logits = model(image_tensor)
    probs = torch.softmax(logits, dim=1)
    predicted_class = torch.argmax(probs, dim=1).item()

class_names = ["Cat", "Dog"]

label = class_names[predicted_class]

plt.figure(figsize=(8, 6))
plt.imshow(image)
plt.axis('off')
plt.title(f'Prediction: {class_names[predicted_class]} ({probs[0, predicted_class].item() * 100:.2f}%)', fontsize=14, fontweight='bold')
plt.tight_layout()
plt.show()
```

Browse... b54472aacb5d6ee99f29c6170fa6725e.JPG

**b54472aacb5d6ee99f29c6170fa6725e.JPG**(image/jpeg) - 279456 bytes, last modified: n/a - 100% done

Saving b54472aacb5d6ee99f29c6170fa6725e.JPG to b54472aacb5d6ee99f29c6170fa6725e.JPG

**Prediction: Dog (88.86% confidence)**



